

Group-MoE: Learned Dispatch to Algebraic Fixed-Function Units

Robb Walters

April 2026 | arXiv Preprint

Abstract

Neural networks that encounter symmetry either enforce it rigidly in every layer or ignore it entirely. We propose GROUP-MoE, a mixture-of-experts architecture where each expert implements a finite group representation in the irreducible representation (irrep) basis, and a learned router detects which symmetry applies. This gives models *algebraic fixed-function units*—analogous to dedicated hardware accelerators—that the network learns to invoke selectively. We introduce the *complement split*, a controlled evaluation methodology that isolates symmetry transfer from memorization, and validate the architecture on synthetic tasks. Our experiments demonstrate: (1) complement transfer from seen to unseen orderings, with the routing architecture providing a robust +3.2pp advantage over plain baselines on S_3 , (2) near-perfect zero-shot compositional generalization from transpositions to 3-cycles ($98.7 \pm 0.8\%$), (3) correct router discrimination between non-nested groups, and (4) zero-degradation compatibility as a transformer FFN replacement. We find that the routing architecture consistently helps, while the fixed irrep matrices provide the strongest advantage on specific seeds but do not reliably outperform learned transforms across seeds. We analyze when algebraic structure outperforms lookup-table memorization and identify the combinatorial scaling regime where group experts should provide their greatest advantage.

1 Introduction

Symmetry is ubiquitous in the data that neural networks process. Images have rotational and reflective symmetries. Molecules have point-group symmetries. Language has permutation symmetries: “Alice likes Bob” and “Bob is liked by Alice” describe the same relationship regardless of entity ordering. The question is not whether symmetry exists, but how a model should handle it.

The field has converged on two approaches, each with significant limitations. *Geometric deep learning* [1, 2] builds equivariance into every layer, ensuring that symmetry transformations of the input produce corresponding transformations of the output. This is correct by construction but rigid: the model cannot mix symmetric and non-symmetric computation, and applying the wrong symmetry is as harmful as applying none. *Standard architectures* (MLPs, transformers) ignore symmetry entirely and hope the right behavior emerges from training data. Our prior work [11] showed that large language models do achieve *functional* equivariance—correct outputs under input permutation—but without developing *structural* equivariance—group representations in their latent space. They memorize the multiplication table rather than learning algebra.

This paper proposes a third path: **group representations as optional expert modules** with learned dispatch. The analogy to hardware design is precise. Just as processors evolved from general-purpose CPUs to include fixed-function accelerators (FFT units, matrix multiply engines, cryptographic co-processors), neural architectures can include *algebraic fixed-function units*—modules hardwired in the irreducible representation basis of a finite group—that are activated only when the router detects the corresponding symmetry. The irrep basis provides the same kind of structural compression that fixed-function hardware provides over general-purpose computation: a $k \times k$ block-diagonal transform instead of a $d \times d$ dense matrix, with exact composition guaranteed by the group homomorphism property.

Our contributions:

1. The GROUP-MoE architecture: group representations in the irrep basis as expert modules, combined with a learned symmetry router and residual blending (Section 2).

2. The *complement split*: a controlled evaluation methodology that isolates symmetry transfer from memorization by ensuring every test example requires generalization from a seen ordering to an unseen one (Section 3).
3. Experimental validation on five synthetic tasks, each answering one question about the architecture (Section 4).
4. Analysis of when algebraic structure outperforms lookup-table memorization, connecting our findings to the broader question of scaling (Section 5).

2 Architecture

2.1 Group Expert Module

The core building block of GROUP-MOE is the *group expert*, which applies a group representation in a learned subspace of the activation space. For a finite group G with irreducible representations of total dimension k , the expert implements:

$$x \xrightarrow{P} z \xrightarrow{R(g)} z' \xrightarrow{Q} x' \quad (1)$$

where $x \in \mathbb{R}^d$ is the input activation, $P \in \mathbb{R}^{k \times d}$ is a learned projection into the symmetry-active subspace, $R(g) \in \mathbb{R}^{k \times k}$ is the irrep matrix for group element g (fixed, not learned), and $Q \in \mathbb{R}^{d \times k}$ is a separately learned injection back to the full space (not constrained to be the pseudo-inverse of P).

The irrep matrix $R(g)$ is block-diagonal by construction:

$$R(g) = \begin{pmatrix} \rho_1(g) & & \\ & \rho_2(g) & \\ & & \ddots \end{pmatrix} \quad (2)$$

where each ρ_i is an irreducible representation. For S_3 , the symmetric group on three elements, the irreps are: trivial (1D), sign (1D), and standard (2D), giving $k = 4$. For a model with $d = 128$, this provides $32 \times$ compression relative to a full $d \times d$ matrix.

The critical property is **exact composition**: $R(g_1)R(g_2) = R(g_1g_2)$ by the group homomorphism. This means the expert’s behavior on composed group elements is determined entirely by its behavior on generators—a property no learned transform can guarantee.

2.2 Symmetry Router

The router is a lightweight MLP that maps each activation to a distribution over routing options: one pass-through option plus one option per group element across all available groups. For a model with groups $G_1, \dots, G_m$:

$$\text{options} = \{\text{pass-through}\} \cup \bigcup_{i=1}^m \{g \mid g \in G_i\} \quad (3)$$

The router uses hard routing (argmax) with soft confidence blending:

$$x' = x + \alpha \cdot (\text{expert}(x, g) - x) \quad (4)$$

where α is the router’s confidence score for the selected option. Pass-through tokens ($\alpha = 0$ or pass-through selected) pass unchanged, ensuring that symmetry is only applied when detected.

What the router learns. The router’s primary role is to decide *whether* to apply a group expert, not to identify the specific permutation relating test to training orderings. The projection P absorbs the element-specific mapping: it learns to project activations into the irrep subspace such that the appropriate block-diagonal transformation $R(g)$ produces the correct output. The router selects among group elements based on the activation pattern, and the confidence α modulates the expert’s contribution. In practice, we observe that the router learns soft operation-level preferences (e.g., higher S_2 routing for commutative operations) rather than hard per-example element identification.

2.3 Load Balancing

Following the Switch Transformer [5], we add an auxiliary load-balancing loss to prevent expert collapse:

$$\mathcal{L}_{\text{balance}} = N \sum_{i=1}^N f_i \cdot p_i \quad (5)$$

where N is the number of routing options, f_i is the fraction of tokens assigned to option i (hard), and p_i is the mean routing probability for option i (soft). This encourages uniform utilization of all options without forcing specific routing patterns.

2.4 Integration

The GROUP-MoE layer is a residual module that handles both (B, d) and (B, L, d) inputs by flattening the sequence dimension internally. This makes it a drop-in replacement for any feedforward block in a transformer or MLP architecture.

Table 1 lists the implemented groups with their irrep dimensions and parameter counts.

Table 1: Implemented group representations.

Group	Order	Irreps	k	Expert params ($d=128$)
$\mathbb{Z}_2$	2	trivial + sign	2	512
$\mathbb{Z}_3$	3	trivial + standard(2D)	3	768
S_2	2	trivial + sign	2	512
S_3	6	trivial + sign + standard(2D)	4	1,024

3 The Complement Split

Standard random train/test splits do not isolate symmetry-specific generalization. A model might generalize from $(a, +, b)$ to $(b, +, a)$ through embedding interpolation rather than by learning commutativity. To measure the group expert’s structural contribution, we need a split where every test example *requires* symmetry-specific transfer.

Definition. For a symmetric function f on n inputs with symmetry group G , the *complement split* assigns exactly one ordering per G -orbit to training and all remaining orderings to test. Formally, for each orbit $\mathcal{O} = \{g \cdot (a_1, \dots, a_n) \mid g \in G\}$, we select one representative $\sigma \in \mathcal{O}$ for training; the remaining $|\mathcal{O}| - 1$ elements go to test.

For S_2 on pairs $\{a, b\}$: train on $(a, +, b)$, test on $(b, +, a)$. Every test example is a reversal of a training example.

For S_3 on triples $\{a, b, c\}$: train on 1 ordering, test on the other 5. The generalization ratio is 1:5—substantially harder than S_2 ’s 1:1.

Composition split (variant): train on identity + transpositions (4 orderings), test on 3-cycles (2 orderings). This tests compositional generalization: the 3-cycle orderings are compositions of transpositions, and the irrep basis computes them by construction ($R(\sigma_1\sigma_2) = R(\sigma_1)R(\sigma_2)$), while a model without group structure must learn them independently.

In all cases, a non-symmetric control operation is split randomly, providing a baseline for the model’s general learning capacity on the same data. Figure 1 illustrates the complement split for S_3 .


4 Experiments

We present five experiments, each answering one question. All use the same base architecture: opaque embeddings (one per number, shared across positions), a concatenation layer, and residual MLP blocks, followed by the group expert layer and a scalar regression head. Unless noted, $d = 128$, 2 MLP blocks, AdamW with learning rate 3×10^{-4} , weight decay 10^{-2} , balance loss $\alpha = 0.01$.

Complement Split for $\{1, 2, 3\}$

TRAIN (1 ordering)

TEST (5 orderings)



All orderings compute the same value: $f(1,2,3) = 1+2+3 = 6$

Group-MoE transfers via $R(g)$; baseline must learn each independently

Figure 1: The complement split for S_3 on the triple $\{1, 2, 3\}$. One ordering (green) goes to training; the remaining five permutations (red) go to test. All orderings compute the same value under the symmetric function. A model with S_3 group structure transfers by construction; a model without it must learn each ordering independently.

4.1 Does complement transfer work? (S_2 , Arithmetic)

Task. $(a, \text{op}, b) \rightarrow \text{result}$, $\text{op} \in \{+, -\}$, $a, b \in [0, 20)$. Complement split on addition pairs: train on $(a, +, b)$, test on $(b, +, a)$. Subtraction split randomly.

Result. GROUP-MoE achieves 48.4% complement accuracy vs. 31.6% for the baseline—a 53% relative improvement. The effect is consistent across 3 random seeds (Table 2).

Table 2: S_2 complement transfer across seeds (num_range=20).

Seed	GROUP-MoE	Baseline	Relative gain
42	48.4%	31.6%	+53%
123	45.8%	27.9%	+64%
7	35.3%	33.2%	+6%

Router behavior. On the complement split, the router routes 70% of addition examples to the S_2 expert vs. 64% of subtraction examples—a consistent preference for the commutative operation. A separate per-pair analysis shows GROUP-MoE exclusively solves $2.5\times$ more complement pairs than the baseline (40 vs. 16 out of 190).

Seed variability. The advantage is consistent in direction (Group-MoE always outperforms) but varies in magnitude: seeds 42 and 123 show strong effects (+53%, +64%), while seed 7 shows a modest +6%. Investigation shows that seed 7 produces weaker router discrimination (the router does not cleanly separate addition from subtraction), suggesting the effect size depends on the router finding a good initialization. The direction is robust; the magnitude depends on router quality.

Balance loss ablation. $\alpha = 0.01$ is the sweet spot. $\alpha = 0$ still produces complement transfer (+30% relative) but with weaker router discrimination. $\alpha = 0.1$ forces excessive routing and hurts performance.

4.2 Is it the group structure or just the routing? (S_3 , Three-Way)

A skeptic could attribute the advantage to the routing architecture rather than the group structure. We add a controlled middle ground: STANDARDMoE, which has the same router, the same project-inject dimensions ($d \rightarrow k \rightarrow d$), the same residual blending—but replaces the fixed irrep matrix $R(g)$ with a *learned* $k \times k$ matrix W . The three models have nearly identical parameter counts ($\sim 337\text{K}$).




Figure 2: Per-pair analysis of S_2 complement transfer (num_range=20, seed=42). **Top row:** absolute prediction error for Group-MoE (left) and baseline (right) on addition test pairs $(b, +, a)$. Green = low error. **Bottom left:** advantage map (blue = Group-MoE has lower error). **Bottom right:** router decisions—dark blue indicates routing to the S_2 expert. The router consistently activates S_2 for addition pairs across the number space.

Task. $(a, b, c, \text{op}) \rightarrow \text{result}$, $\text{op} \in \{\text{sum}, \text{nonsym}\}$, num_range=15. Two split modes: complement (1 ordering trains, 5 test) and composition (identity + transpositions train, 3-cycles test).

Table 3: Three-way comparison on S_3 ternary task (num_range=15). Mean $\pm$ std over 5 seeds.

Model	Complement (1 $\rightarrow$ 5)	Composition (4 $\rightarrow$ 2)
GROUP-MoE (irrep $R(g)$)	89.2 $\pm$ 3.1%	98.7 $\pm$ 0.8%
StandardMoE (learned W)	90.2 $\pm$ 2.0%	98.6 $\pm$ 0.8%
Baseline (no expert)	87.0 $\pm$ 1.6%	97.4 $\pm$ 1.1%

Decomposition. Across 5 seeds, the routing architecture provides a robust +3.2pp advantage over the baseline on complement split (StandardMoE vs. Baseline). However, the fixed irrep matrices do not reliably outperform learned transforms: GROUP-MoE averages -1.0pp relative to StandardMoE on complement, with high seed variance ($\sigma = 3.1\%$). On composition, GROUP-MoE and StandardMoE are essentially tied (98.7% vs. 98.6%), both beating the baseline by $\sim 1.3\text{pp}$ (Figure 3).

Interpretation. The routing architecture—project, transform, inject with learned dispatch—is the primary source of complement transfer at this scale. The fixed irrep structure provides its strongest advantage on favorable seeds (up to +3.6pp over StandardMoE) but hurts on unfavorable ones (-5.5pp on seed 7),




Figure 3: Three-way comparison on S_3 ternary task (num_range=15). **Left:** complement split (1 ordering trains, 5 test). **Right:** composition split (transpositions train, 3-cycles test). Red annotations show the decomposition: +3.3pp from group structure, +1.7pp from routing architecture.

suggesting the router’s ability to exploit the irrep basis depends on initialization. The irrep basis guarantees exact composition—a property no learned transform can match—but this guarantee does not translate to consistent empirical gains in our small-scale synthetic setting.

4.3 Does irrep composition generalize? (Transpositions $\rightarrow$ 3-Cycles)

The strongest theoretical prediction of group representations: since $R(\sigma_1\sigma_2) = R(\sigma_1)R(\sigma_2)$, the expert should work for 3-cycles without ever training on them.

Setup. Composition split: train on identity + 3 transpositions per triple (generators of S_3), test only on the 2 three-cycle orderings (compositions).

Result. Across 5 seeds, GROUP-MoE achieves $98.7 \pm 0.8\%$ on 3-cycle orderings never seen during training. The projection P , learned from transposition examples, generalizes because the irrep subspace is shared across all group elements. At num_range=10, the router achieves perfect discrimination: 100% S_3 routing for the symmetric operation vs. 74% for the non-symmetric control.

4.4 Can the router discriminate between groups?

Nested groups ($S_2 + S_3$). With both S_2 and S_3 experts available, the router routes $\sim 70\%$ of all operations to S_3 regardless of operation type. This is *rational*: since $S_2 \subset S_3$, every S_2 transformation is also an S_3 transformation. The router discovers the subgroup relationship without supervision.

Non-nested groups ($\mathbb{Z}_2 + \mathbb{Z}_3$). With genuinely disparate groups (neither is a subgroup of the other), the router correctly dispatches $\mathbb{Z}_2$ operations to the $\mathbb{Z}_2$ expert at 76%. The non-symmetric control routes primarily to pass-through (15%) or uses the larger expert as general-purpose compute.

This confirms that the router learns which algebraic unit to invoke when the units have non-overlapping capabilities (Figure 4).

4.5 Transformer compatibility

We also tested GROUP-MoE as a drop-in FFN replacement in a 4-layer transformer encoder ($d = 256$, 4 heads), with each of (a, op, b, c) as a separate token. All three models reach 100% complement accuracy: self-attention on 4 tokens provides functional equivariance natively, making the group expert’s marginal




Figure 4: Router behavior with multiple group experts. **Left:** nested groups ($S_2 \subset S_3$)—the router rationally routes everything to the more expressive S_3 expert. **Right:** non-nested groups (Z_2, Z_3)—the router correctly dispatches Z_2 operations to the Z_2 expert at 76%.

contribution negligible for this short-sequence task. This confirms architectural compatibility—GROUP-MOE introduces zero degradation—while illustrating that the group expert’s advantage requires tasks where attention alone cannot resolve the symmetry (e.g., symmetry among a subset of tokens in a longer sequence).

Reproducibility. All experiments use PyTorch on Apple Silicon (MPS). Code, datasets, and training scripts are available at <https://github.com/rwalters/group-moe>. Each experiment runs in under 10 minutes on a single machine.

5 When Does Algebraic Structure Beat Memorization?

Our experiments reveal a consistent pattern that we frame as the **lookup table (LUT) vs. algebra** tradeoff.

Memorization wins when: (1) the function decomposes into per-element contributions (e.g., sum = per-number values), so the effective table size is $O(n)$ not $O(n^k)$; (2) the model has enough capacity relative to the table size; (3) attention provides permutation mixing natively; (4) enough orderings are observed that interpolation suffices.

Algebraic structure helps when: (1) the complement split isolates symmetry transfer from memorization; (2) the model’s backbone does not already handle permutation mixing; (3) the effective lookup table has gaps that the group structure can fill; (4) the irrep composition property enables zero-shot generalization to unseen composed elements.

A scaling prediction (untested). The group S_n has $n!$ elements. For S_{10} , this is 3.6 million; for S_{20} , 2.4×10^{18} . No model memorizes 10^{18} entries, but the irrep basis for S_{20} remains structured and compressible. Our synthetic experiments operate in the small-table regime where memorization is competitive—our scaling studies (Appendix C) confirm that increasing the input range does not produce a crossover because the sum function decomposes into per-element contributions. We predict that the advantage of GROUP-MOE will manifest in the *large-table regime*—large groups, long sequences, combinatorial state spaces—where the factorial growth of the lookup table defeats memorization. Validating this prediction requires tasks with larger effective symmetry groups and is an important direction for future work.

The developmental analogy is apt: humans learn multiplication tables by rote before understanding algebra. For small tables, rote memorization is strictly easier. Algebra pays off when the table is too large to memorize—and in the limit, it is the *only* strategy that works.

6 Related Work

Geometric deep learning. Cohen & Welling [2] introduced group equivariant convolutions; Bronstein et al. [1] unified CNNs, GNNs, and transformers under the equivariance framework; Weiler & Cesa [3] developed general steerable CNNs using irrep decomposition. All enforce symmetry in every layer. GROUP-MoE uses the same irrep machinery but makes it optional via routing.

Mixture of experts. Shazeer et al. [4] proposed sparse MoE with gating; Fedus et al. [5] simplified to top-1 routing with load balancing. Our router and balance loss follow this design, but our experts have algebraic structure rather than being generic MLPs. Kang et al. [6] propose “Mixture of Group Experts” using group sparse regularization for expert diversity; despite the shared name, their “groups” refer to expert grouping, not group theory.

Learned group representations. MatrixNet [7] learns matrix representations of group elements from data, achieving compositional generalization via learned generator representations. Our approach is complementary: we fix the irrep structure (guaranteeing exact composition) and learn the dispatch (routing).

Symmetry detection. Dehmamy et al. [8] discover continuous symmetries via Lie algebra convolution; Benton et al. [9] meta-learn invariances. These discover symmetries post-hoc. GROUP-MoE provides exact group structure as architectural options with learned selection.

Compositional generalization. Gordon et al. [10] hypothesize language compositionality is group-equivariance and build equivariant seq2seq models. We share the hypothesis but provide group structure as optional expert modules rather than mandatory architectural constraints.

Selective and partial equivariance. Several recent works allow models to modulate the degree of equivariance. PEnGUIN [12] introduces learnable symmetry scores controlling per-layer equivariance in GNNs. Puny et al. [13] propose frame averaging to break symmetry in equivariant networks when needed. These approaches start from full equivariance and learn when to relax it; GROUP-MoE takes the dual approach, starting from no equivariance and learning when to impose it.

Functional vs. structural equivariance. Walters [11] (“Ordering Is Not Invariant”) showed LLMs achieve functional equivariance without structural equivariance. GROUP-MoE is the architectural response: providing structural equivariance as an option.

7 Discussion and Future Work

The ASIC analogy. Group experts are cognitive arithmetic units—domain-specific algebraic accelerators hardwired in the irrep basis, dispatched by a learned router. The analogy to hardware evolution is not merely metaphorical: the irrep basis provides the same structural compression that fixed-function hardware provides over general-purpose computation, and the router serves the same dispatch role as a processor’s instruction decoder.

Scaling to combinatorial regimes. Our synthetic experiments validate the mechanism but operate where memorization is competitive. The key prediction: for S_n with large n , the $n!$ growth of the lookup table will defeat memorization while the irrep basis remains tractable. Testing this requires implementing larger permutation groups or finding natural tasks with large effective symmetry groups.

Language modeling. Natural language contains implicit symmetry: entity permutations, fact reorderings, syntactic transforms. These create effective symmetry groups whose “multiplication tables” are too large for rote memorization. Inserting GROUP-MoE into language models and testing with entity permutation probes [11] is the natural next step.

Symmetry discovery. Currently the groups are specified a priori. A compelling extension: learn which groups are useful from data, by treating the irrep structure as differentiable or using a library of candidate groups with learned selection.

Approximate symmetry. Real data has approximate, not exact, symmetry. Soft irrep decomposition with learned tolerance could extend GROUP-MoE to noisy settings.

8 Conclusion

GROUP-MOE demonstrates that neural networks can learn to dispatch to algebraic fixed-function units. The complement split methodology cleanly isolates symmetry transfer from memorization. The three-way comparison (5 seeds) shows that the routing architecture is the primary driver of complement transfer (+3.2pp), while the fixed irrep matrices provide strong gains on favorable seeds but do not consistently outperform learned transforms. Compositional generalization confirms the theoretical payoff: irrep composition enables near-perfect zero-shot transfer to composed elements ($98.7 \pm 0.8\%$). The router learns group-theoretic relationships—subgroup containment, non-nested discrimination—without supervision.

These are proof-of-concept results on synthetic tasks where memorization is competitive. The architecture’s value will compound in regimes where the effective symmetry group is too large to memorize—exactly the combinatorial territory where current models struggle. GROUP-MOE provides the algebraic tools; identifying the right large-scale application is future work.

References

- [1] M. M. Bronstein, J. Bruna, T. Cohen, and P. Veličković. Geometric deep learning: Grids, groups, graphs, geodesics, and gauges. *arXiv:2104.13478*, 2021.
- [2] T. S. Cohen and M. Welling. Group equivariant convolutional networks. In *ICML*, 2016.
- [3] M. Weiler and G. Cesa. General E(2)-equivariant steerable CNNs. In *NeurIPS*, 2019.
- [4] N. Shazeer, A. Mirhoseini, K. Maziarz, A. Davis, Q. Le, G. Hinton, and J. Dean. Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. In *ICLR*, 2017.
- [5] W. Fedus, B. Zoph, and N. Shazeer. Switch Transformers: Scaling to trillion parameter models with simple and efficient sparsity. *JMLR*, 23(120):1–40, 2022.
- [6] L. Kang, J. Li, M. Tian, and H. Huang. Mixture of group experts for learning invariant representations. *arXiv:2504.09265*, 2025.
- [7] L. Laird, C. Hsu, A. Bapat, and R. Walters. MatrixNet: Learning over symmetry groups using learned group representations. In *NeurIPS*, 2024.
- [8] N. Dehmamy, R. Walters, Y. Liu, D. Wang, and R. Yu. Automatic symmetry discovery with Lie algebra convolutional network. In *NeurIPS*, 2021.
- [9] G. Benton, M. Finzi, P. Izmailov, and A. G. Wilson. Learning invariances in neural networks from training data. In *NeurIPS*, 2020.
- [10] J. Gordon, D. Lopez-Paz, M. Baroni, and D. Bouchacourt. Permutation equivariant models for compositional generalization in language. In *ICLR*, 2020.
- [11] R. Walters. Ordering is not invariant: Probing structural equivariance in large language models. 2025.
- [12] J. McClellan, G. Brothers, F. Huang, and P. Tokekar. PEnGUIN: Partially equivariant graph neural networks for sample efficient MARL. *Reinforcement Learning Journal*, 6:2637–2651, 2025.
- [13] O. Puny, M. Atzmon, E. J. Smith, I. Misra, A. Grover, H. Ben-Hamu, and Y. Lipman. Frame averaging for invariant and equivariant network design. In *ICLR*, 2022.

A Group Representation Details

All irrep matrices are verified to satisfy the group composition property: $R(g_1)R(g_2) = R(g_1g_2)$ for all pairs (g_1, g_2) . Verification is automated in the test suite (`tests/test_groups.py`).

S_3 irrep matrices. The six elements of S_3 are: identity e , three transpositions (01) , (12) , (02) , and two 3-cycles (012) , (021) . In the irrep basis $[\text{trivial} \oplus \text{sign} \oplus \text{standard}]$, each element maps to a 4×4 block-diagonal matrix. For example, the transposition (01) maps to $\text{diag}(1, -1) \oplus \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}$, and the 3-cycle (012) maps to $\text{diag}(1, 1) \oplus \begin{pmatrix} \cos(2\pi/3) & -\sin(2\pi/3) \\ \sin(2\pi/3) & \cos(2\pi/3) \end{pmatrix}$.

Full multiplication tables, irrep matrices, and composition verification code are available in the codebase (`src/groups/representations.py`).

B Hyperparameter Sensitivity

Balance loss α . Tested at $\alpha \in \{0, 0.01, 0.1\}$ on S_2 arithmetic complement split. $\alpha = 0.01$ is the sweet spot: $\alpha = 0$ provides complement transfer without router discrimination; $\alpha = 0.1$ forces excessive routing.

Model dimension d . Tested at $d \in \{32, 64, 128, 256\}$. Results are qualitatively consistent; the group structure’s advantage is proportionally similar across scales.

Number range. At very small `num_range` (≤ 10), embeddings lack sufficient training context and all models underperform. At `num_range` ≥ 15 , the sum function is learnable and the complement split provides clear signal.

C Scaling Analysis

We conducted two scaling studies to determine whether the group advantage grows with problem size.

num_range sweep (10–50, $d = 128$ and $d = 32$): all models converge to 100% complement accuracy by `num_range` = 25. The sum function decomposes into per-number values, so the effective table size grows as $O(n)$ rather than $O(n^3)$. This means `num_range` scaling does not create the combinatorial challenge needed to separate the models.

Coverage sweep (1–4 orderings per triple, `num_range` = 15): differentiation exists only at $k = 1$ (17% coverage, 1:5 generalization). By $k = 2$ (33%), all models converge to >99%. The jump from $k = 1$ to $k = 2$ is the critical transition where even a single additional ordering provides enough context for memorization to succeed.

These results confirm that the sum function, while useful for validating the architecture, does not create a combinatorial scaling challenge. The right scaling axis is group order ($n!$ growth), not input range (n growth).